

Part 01 :

1. Consider this dataset from kaggle. (Download the dataset from following link : <https://www.kaggle.com/shrutimechlearn/step-by-step-kmeans-explained-in-detail/data>) and answer the following questions :

Importing the necessary lib

```
In [1]: import warnings
warnings.filterwarnings('ignore')

In [2]: import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn import metrics
from sklearn.metrics import accuracy_score, classification_report

In [3]: import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set(style="whitegrid", color_codes=True, palette="dark" )
```

Load Dataset

```
In [4]: mall_customers = pd.read_csv("Mall_Customers.csv")
```

Exploratory Data Analysis (EDA)

```
In [5]: mall_customers.sample(10)
```

```
Out[5]:
```

	CustomerID	Genre	Age	Annual_Income_(k\$)	Spending_Score
152	153	Female	44	78	20
61	62	Male	19	46	55
49	50	Female	31	40	42
196	197	Female	45	126	28
135	136	Female	29	73	88
77	78	Male	40	54	48
18	19	Male	52	23	29
78	79	Female	23	54	52
9	10	Female	30	19	72
33	34	Male	18	33	92

```
In [6]: mall_customers.dtypes
```

```
Out[6]:
```

CustomerID	int64
Genre	object
Age	int64
Annual_Income_(k\$)	int64
Spending_Score	int64
dtype:	object

```
In [7]: mall_customers.shape
```

```
Out[7]: (200, 5)
```

```
In [8]: duplicate_rows = mall_customers[mall_customers.duplicated()]

print("Duplicate Rows:")
duplicate_rows
```

Duplicate Rows:

```
Out[8]:
```

	CustomerID	Genre	Age	Annual_Income_(k\$)	Spending_Score
--	------------	-------	-----	---------------------	----------------

```
In [9]: mall_customers.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Genre                 200 non-null   object
2   Age                   200 non-null   int64
3   Annual_Income_(k$)    200 non-null   int64
4   Spending_Score        200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
In [10]: mall_customers.isnull().sum()
```

```
Out[10]:
```

CustomerID	0
Genre	0
Age	0
Annual_Income_(k\$)	0
Spending_Score	0
dtype:	int64


```
In [11]: print(mall_customers.groupby('Genre').size())

sns.countplot(x='Genre', data=mall_customers)
plt.title("Distribution of Customers by Gender")
plt.show()
```

Genre

Female	112
Male	88

dtype: int64



```
In [12]: mall_customers.describe().T
```

Out[12]:

	count	mean	std	min	25%	50%	75%	max
CustomerID	200.0	100.50	57.879185	1.0	50.75	100.5	150.25	200.0
Age	200.0	38.85	13.969007	18.0	28.75	36.0	49.00	70.0
Annual_Income_(k\$)	200.0	60.56	26.264721	15.0	41.50	61.5	78.00	137.0
Spending_Score	200.0	50.20	25.823522	1.0	34.75	50.0	73.00	99.0

Checking skewness

```
In [13]: numeric_skewness = mall_customers.select_dtypes(include=[np.number]).skew()
print("Skewness:")
print(numeric_skewness)
```

Skewness:
CustomerID 0.000000
Age 0.485569
Annual_Income_(k\$) 0.321843
Spending_Score -0.047220
dtype: float64

Applying Label Encoding to "Genre" column.

```
In [14]: mall_customers_encoded = mall_customers.copy()
label_encoder = LabelEncoder()
mall_customers_encoded['Genre'] = label_encoder.fit_transform(mall_customers_encoded['Genre'])
mall_customers_encoded
```

Out[14]:

	CustomerID	Genre	Age	Annual_Income_(k\$)	Spending_Score
0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6
3	4	0	23	16	77
4	5	0	31	17	40
...
195	196	0	35	120	79
196	197	0	45	126	28
197	198	1	32	126	74
198	199	1	32	137	18
199	200	1	30	137	83

200 rows x 5 columns

1.1 Perform k-means clustering over this dataset using Manhattan distance as the distance-measure. (10 Points)

Def function for manhattan_distance measure

```
In [15]: def manhattan_distance(point1, point2):
        """
        Calculate the Manhattan distance between two points.

        Parameters:
        - point1: Tuple or list representing the coordinates of the first point.
        - point2: Tuple or list representing the coordinates of the second point.

        Returns:
        - The Manhattan distance between the two points.
        """
        distance = sum(abs(x - y) for x, y in zip(point1, point2))
        return distance
```

k-means clustering with Manhattan distance

```
In [16]: features = mall_customers[['Age', 'Annual_Income_(k$)', 'Spending_Score']].values

k = 5
kmeans = KMeans(n_clusters=k, init='k-means++', algorithm='elkan', n_init=10, random_state=34)

kmeans.fit(features)

cluster_centers = kmeans.cluster_centers_
cluster_labels = kmeans.labels_

manhattan_distances = manhattan_distance(features, cluster_centers)

print("Cluster Centers:")
cluster_centers_df = pd.DataFrame(cluster_centers, columns=['Age', 'Annual_Income_(k$)', 'Spending_Score'])
print(cluster_centers_df, '\n')
```

```
print("Manhattan Distances to Cluster Centers:", manhattan_distances)
```

Cluster Centers:

	Age	Annual_Income_(k\$)	Spending_Score
0	43.282051	55.025641	49.692308
1	32.692308	86.538462	82.128205
2	25.521739	26.304348	78.565217
3	45.217391	26.304348	20.913043
4	40.324324	87.432432	18.189189

Manhattan Distances to Cluster Centers: [73.03781373 202.60523065 162.28349754]

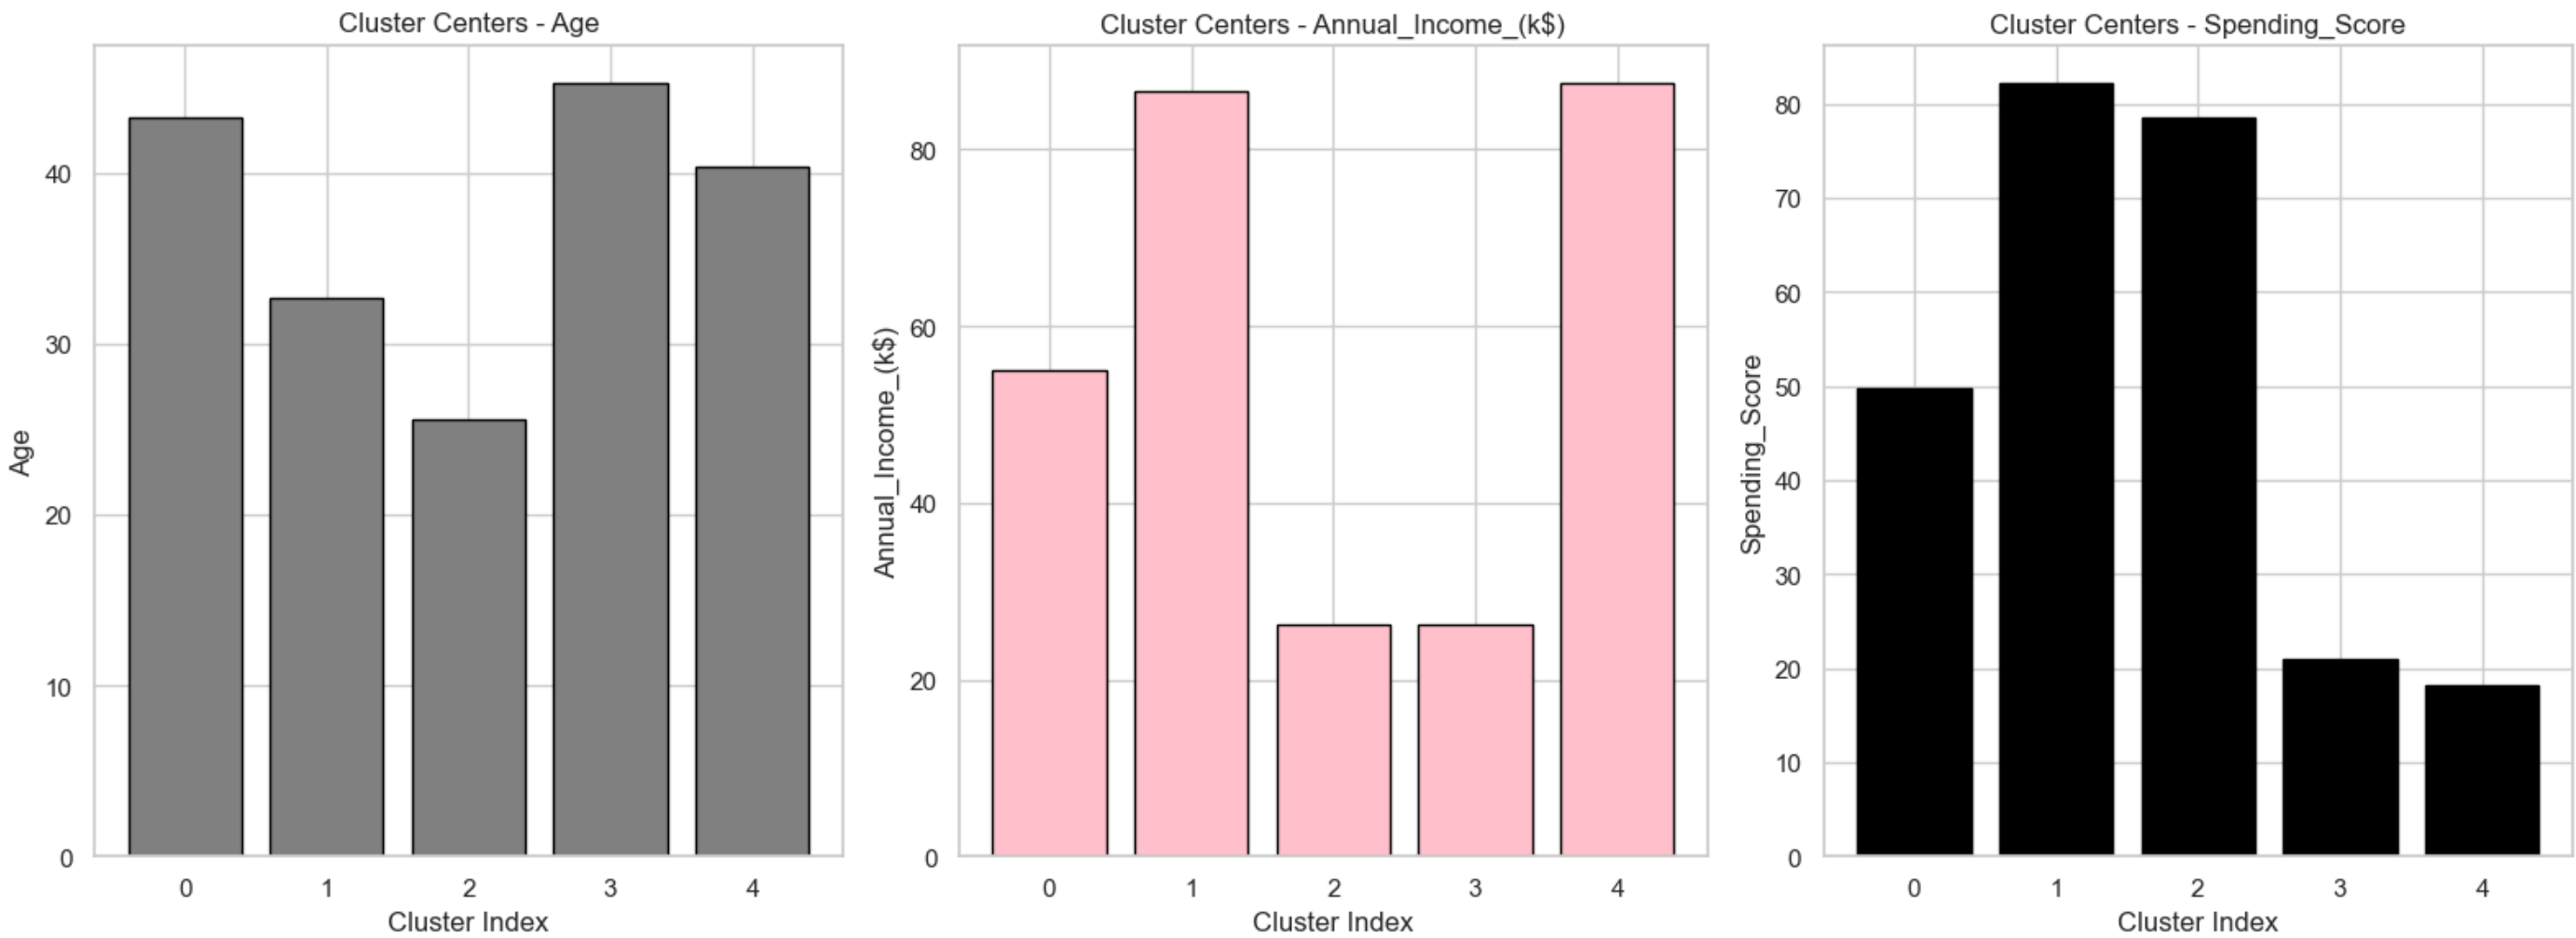
Visualization of Cluster Center output using histogram

```
In [17]: plt.figure(figsize=(16, 6))

colors = ['gray', 'pink', 'black']

for i, feature in enumerate(cluster_centers_df.columns):
    plt.subplot(1, 3, i+1)
    plt.bar(range(len(cluster_centers_df)), cluster_centers_df[feature], color=colors[i], edgecolor = 'black')
    plt.xlabel('Cluster Index')
    plt.ylabel(feature)
    plt.title(f'Cluster Centers - {feature}')
    plt.xticks(range(len(cluster_centers_df)))
    plt.tight_layout()

plt.show()
```




Visualization for Manhattan Distances to Cluster Centers using histogram and box-plot

```
In [101... plt.figure(figsize=(16, 6))

plt.subplot(1, 2, 1)
plt.bar(range(len(manhattan_distances)), manhattan_distances, color='Brown', edgecolor = 'black')
plt.xlabel('Cluster Index')
plt.ylabel('Manhattan Distance')
plt.title('Manhattan Distances to Cluster Centers')
plt.xticks(range(len(manhattan_distances)))

plt.subplot(1, 2, 2)
plt.boxplot(manhattan_distances)
plt.xlabel('All Clusters')
plt.ylabel('Manhattan Distance')
plt.title('Manhattan Distances to Cluster Centers')

plt.tight_layout()
plt.show()
```



Visualization of the clusters to check if they have a distinct distinction in the feature space.


```
In [96]: pca = PCA(n_components=2)
data_pca = pca.fit_transform(features)

tsne = TSNE(n_components=2, random_state=42)
data_tsne = tsne.fit_transform(features)

plt.figure(figsize=(20, 8))
plt.subplot(1, 2, 1)
for label in set(cluster_labels):
    plt.scatter(data_pca[cluster_labels == label, 0], data_pca[cluster_labels == label, 1], label=f'Cluster {label}')
plt.title('PCA Visualization of Clusters')
plt.xlabel('Principal Component 1')
```

```
plt.ylabel('Principal Component 2')
plt.legend()

plt.subplot(1, 2, 2)
for label in set(cluster_labels):
    plt.scatter(data_tsne[cluster_labels == label, 0], data_tsne[cluster_labels == label, 1], label=f'Cluster {label}')
plt.title('t-SNE Visualization of Clusters')
plt.xlabel('t-SNE Component 1')
plt.ylabel('t-SNE Component 2')
plt.legend()
plt.tight_layout()
plt.show()
```



The above output, displays the scatter plot for both PCA and t-SNE to identify clusters within the data.

The clusters appear to be distinct and well-separated in the feature space, implying that the data points within each cluster are similar to one another while differentiating from data points in other clusters. This is a good indicator since it shows that the clustering algorithm successfully grouped comparable data points together.

1.2 After performing k-means clustering, extract the groups or clusters and add a separate column in your dataset as ‘Labels’ and fill it with cluster number assigned by k-means algorithm. (5 Points)

Adding the 'Lables' columns to Dataset.

```
In [20]: mall_customers_with_labels = pd.DataFrame(mall_customers_encoded, columns=mall_customers_encoded.columns)
mall_customers_with_labels['Labels'] = cluster_labels
mall_customers_with_labels
```

Out[20]:

	CustomerID	Genre	Age	Annual_Income_(k\$)	Spending_Score	Labels
0	1	1	19	15	39	3
1	2	1	21	15	81	2
2	3	0	20	16	6	3
3	4	0	23	16	77	2
4	5	0	31	17	40	3
...
195	196	0	35	120	79	1
196	197	0	45	126	28	4
197	198	1	32	126	74	1
198	199	1	32	137	18	4
199	200	1	30	137	83	1

200 rows x 6 columns

```
In [21]: print(mall_customers_with_labels.groupby('Labels').size())
```

```
Labels
0    78
1    39
2    23
3    23
4    37
dtype: int64
```

1.3 Now, you should be ready with your labeled dataset. Perform standard classification task using logistic regression, decision trees, random forest, and Naive Bayes algorithm. (25 Points)

```
In [22]: X = mall_customers_with_labels.drop(columns=['Labels'])
y = mall_customers_with_labels['Labels']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=98, stratify=y)

classifiers = {
    "Logistic Regression": LogisticRegression(),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "Naive Bayes": GaussianNB()
}

classifiers_accuracys = {}
for name, clf in classifiers.items():

    clf.fit(X_train, y_train)

    y_pred = clf.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)

    classifiers_accuracys[name] = accuracy
```



```
for name, accuracy in classifiers_accuracys.items():
    print(f"{name}: Accuracy - {(accuracy)*100} %")
```

Logistic Regression: Accuracy - 88.75 %
Decision Tree: Accuracy - 96.25 %
Random Forest: Accuracy - 100.0 %
Naive Bayes: Accuracy - 98.75 %

In [53]: sorted_classifiers_accuracys = sorted(classifiers_accuracys.items(), key=lambda x: x[1], reverse=True)

```
for name, accuracy in sorted_classifiers_accuracys:
    print(f"{name}: {round(accuracy, 4)*100} %")
```

Random Forest: 100.0 %
Naive Bayes: 98.75 %
Decision Tree: 96.25 %
Logistic Regression: 88.75 %

1. Classification Results:

- Random Forest: Accuracy - 100.0%
- Naive Bayes: Accuracy - 98.75%
- Decision Tree: Accuracy - 96.25%
- Logistic Regression: Accuracy - 89.75%

All classification algorithms achieved high accuracies, indicating that they were able to successfully classify the data points into their respective clusters.

1.4 Compare the performance of these various supervised learning algorithm and comment on the homogeneity of clusters, like is the clusters or groups are making sense or not ? (10 Points)

1. Homogeneity of Clusters:

- The classification algorithms' high accuracy ratings indicate that the clusters are somewhat homogeneous.
- High Accuracy: A high accuracy score implies that the supervised learning algorithms accurately classified the data points into their appropriate classes with high confidence. This implies that the data points within each class are largely uniform and distinguishable from those in other classes.
 - Impact on Cluster Homogeneity: While the accuracy of supervised learning methods does not directly assess cluster homogeneity, it does indirectly indicate cluster separation and coherence. Higher accuracy indicates that the clusters or groupings identified by the clustering algorithm make sense and are internally consistent.
- Potential Insights: If the supervised learning algorithms perform well, it means that the clusters identified by the clustering algorithm are significant and represent separate groups in the data. In contrast, low accuracy may signal that the clusters are less uniform and require more refinement.

In conclusion, the high accuracy ratings obtained from the supervised learning methods indicate that the clusters produced by the clustering algorithm are most likely meaningful and internally homogeneous. Otherwise, if classification performance is low or inconsistent, it suggests that the clusters might not be well-separated or meaningful.

Part 02 :

2. Consider the breast_cancer dataset given in the sklearn library and answer the following questions.

Importing the necessary lib

In [24]: from sklearn.datasets import load_breast_cancer
from sklearn.preprocessing import StandardScaler
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

2.1 Import the breast_cancer dataset from sklearn.datasets library. (5 Points)

In [25]: breast_cancer = load_breast_cancer()

Convert to pandas DataFrame

In [26]: breast_cancer_df = pd.DataFrame(breast_cancer.data, columns=breast_cancer.feature_names)
breast_cancer_df['target'] = breast_cancer.target

Exploratory Data Analysis (EDA)

In [27]: pd.set_option('display.max_columns', None)
breast_cancer_df.sample(10)

Out[27]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	radius error	texture error	perimeter error	area error	smoothness error	compactness error	concavity error	concave points error	symmetry error	fractal dimension
128	15.100	16.39	99.58	674.5	0.11500	0.18070	0.11380	0.08534	0.2001	0.06467	0.4309	1.0680	2.796	39.84	0.009006	0.041850	0.03204	0.022580	0.02353	0.0045
486	14.640	16.85	94.21	666.0	0.08641	0.06698	0.05192	0.02791	0.1409	0.05355	0.2204	1.0060	1.471	19.98	0.003535	0.013930	0.01800	0.006144	0.01254	0.0010
369	22.010	21.90	147.20	1482.0	0.10630	0.19540	0.24480	0.15010	0.1824	0.06140	1.0080	0.6999	7.561	130.20	0.003978	0.028210	0.03576	0.014710	0.01518	0.0037
18	19.810	22.15	130.00	1260.0	0.09831	0.10270	0.14790	0.09498	0.1582	0.05395	0.7582	1.0170	5.865	112.40	0.006494	0.018930	0.03391	0.015210	0.01356	0.0015
117	14.870	16.67	98.64	682.5	0.11620	0.16490	0.16900	0.08923	0.2157	0.06768	0.4266	0.9489	2.989	41.18	0.006985	0.025630	0.03011	0.012710	0.01602	0.0038
441	17.270	25.42	112.40	928.8	0.08331	0.11090	0.12040	0.05736	0.1467	0.05407	0.5100	1.6790	3.283	58.38	0.008109	0.043080	0.04942	0.017420	0.01594	0.0037
551	11.130	22.44	71.49	378.4	0.09566	0.08194	0.04824	0.02257	0.2030	0.06552	0.2800	1.4670	1.994	17.85	0.003495	0.030510	0.03445	0.010240	0.02912	0.0047
41	10.950	21.35	71.90	371.1	0.12270	0.12180	0.10440	0.05669	0.1895	0.06870	0.2366	1.4280	1.822	16.97	0.008064	0.017640	0.02595	0.010370	0.01357	0.0030
359	9.436	18.32	59.82	278.6	0.10090	0.05956	0.02710	0.01406	0.1506	0.06959	0.5079	1.2470	3.267	30.48	0.006836	0.008982	0.02348	0.006565	0.01942	0.0020
287	12.890	13.12	81.89	515.9	0.06955	0.03729	0.02260	0.01171	0.1337	0.05581	0.1532	0.4690	1.115	12.68	0.004731	0.013450	0.01652	0.005905	0.01619	0.0020

In [28]: breast_cancer_df.dtypes

```
Out [28]: mean radius          float64
mean texture          float64
mean perimeter        float64
mean area             float64
mean smoothness       float64
mean compactness      float64
mean concavity         float64
mean concave points   float64
mean symmetry         float64
mean fractal dimension float64
radius error          float64
texture error         float64
perimeter error       float64
area error            float64
smoothness error      float64
compactness error     float64
concavity error       float64
concave points error  float64
symmetry error        float64
fractal dimension error float64
worst radius          float64
worst texture         float64
worst perimeter       float64
worst area            float64
worst smoothness      float64
worst compactness     float64
worst concavity       float64
worst concave points  float64
worst symmetry        float64
worst fractal dimension float64
target               int64
dtype: object
```

```
In [29]: breast_cancer_df.shape
```

```
Out [29]: (569, 31)
```

```
In [30]: duplicate_rows = breast_cancer_df[breast_cancer_df.duplicated()]
```

```
if not duplicate_rows.empty:
    print("Duplicate rows found. Details:")
    print(duplicate_rows)
else:
    print("No duplicate rows found.")
```

No duplicate rows found.

```
In [31]: breast_cancer_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   mean radius                          569 non-null    float64
1   mean texture                        569 non-null    float64
2   mean perimeter                      569 non-null    float64
3   mean area                          569 non-null    float64
4   mean smoothness                     569 non-null    float64
5   mean compactness                    569 non-null    float64
6   mean concavity                      569 non-null    float64
7   mean concave points                 569 non-null    float64
8   mean symmetry                      569 non-null    float64
9   mean fractal dimension              569 non-null    float64
10  radius error                        569 non-null    float64
11  texture error                      569 non-null    float64
12  perimeter error                    569 non-null    float64
13  area error                         569 non-null    float64
14  smoothness error                   569 non-null    float64
15  compactness error                  569 non-null    float64
16  concavity error                    569 non-null    float64
17  concave points error               569 non-null    float64
18  symmetry error                     569 non-null    float64
19  fractal dimension error            569 non-null    float64
20  worst radius                       569 non-null    float64
21  worst texture                      569 non-null    float64
22  worst perimeter                    569 non-null    float64
23  worst area                         569 non-null    float64
24  worst smoothness                   569 non-null    float64
25  worst compactness                  569 non-null    float64
26  worst concavity                    569 non-null    float64
27  worst concave points               569 non-null    float64
28  worst symmetry                     569 non-null    float64
29  worst fractal dimension            569 non-null    float64
30  target                             569 non-null    int64
dtypes: float64(30), int64(1)
memory usage: 137.9 KB
```


```
In [32]: null_counts = breast_cancer_df.isnull().sum()
if null_counts.any():
    print("Null values present. Details:")
    print(null_counts)
else:
    print("No null values present.")
```

No null values present.

```
In [33]: target_counts = breast_cancer_df.groupby('target').size()
```

```
target_counts.plot(kind='bar', title='Count of Target Values', xlabel='Target', ylabel='Count')

plt.show()
```



```
In [34]: breast_cancer_df.describe().T
```

Out [34]:

	count	mean	std	min	25%	50%	75%	max
mean radius	569.0	14.127292	3.524049	6.981000	11.700000	13.370000	15.780000	28.11000
mean texture	569.0	19.289649	4.301036	9.710000	16.170000	18.840000	21.800000	39.28000
mean perimeter	569.0	91.969033	24.298981	43.790000	75.170000	86.240000	104.100000	188.50000
mean area	569.0	654.889104	351.914129	143.500000	420.300000	551.100000	782.700000	2501.00000
mean smoothness	569.0	0.096360	0.014064	0.052630	0.086370	0.095870	0.105300	0.16340
mean compactness	569.0	0.104341	0.052813	0.019380	0.064920	0.092630	0.130400	0.34540
mean concavity	569.0	0.088799	0.079720	0.000000	0.029560	0.061540	0.130700	0.42680
mean concave points	569.0	0.048919	0.038803	0.000000	0.020310	0.033500	0.074000	0.20120
mean symmetry	569.0	0.181162	0.027414	0.106000	0.161900	0.179200	0.195700	0.30400
mean fractal dimension	569.0	0.062798	0.007060	0.049960	0.057700	0.061540	0.066120	0.09744
radius error	569.0	0.405172	0.277313	0.111500	0.232400	0.324200	0.478900	2.87300
texture error	569.0	1.216853	0.551648	0.360200	0.833900	1.108000	1.474000	4.88500
perimeter error	569.0	2.866059	2.021855	0.757000	1.606000	2.287000	3.357000	21.98000
area error	569.0	40.337079	45.491006	6.802000	17.850000	24.530000	45.190000	542.20000
smoothness error	569.0	0.007041	0.003003	0.001713	0.005169	0.006380	0.008146	0.03113
compactness error	569.0	0.025478	0.017908	0.002252	0.013080	0.020450	0.032450	0.13540
concavity error	569.0	0.031894	0.030186	0.000000	0.015090	0.025890	0.042050	0.39600
concave points error	569.0	0.011796	0.006170	0.000000	0.007638	0.010930	0.014710	0.05279
symmetry error	569.0	0.020542	0.008266	0.007882	0.015160	0.018730	0.023480	0.07895
fractal dimension error	569.0	0.003795	0.002646	0.000895	0.002248	0.003187	0.004558	0.02984
worst radius	569.0	16.269190	4.833242	7.930000	13.010000	14.970000	18.790000	36.04000
worst texture	569.0	25.677223	6.146258	12.020000	21.080000	25.410000	29.720000	49.54000
worst perimeter	569.0	107.261213	33.602542	50.410000	84.110000	97.660000	125.400000	251.20000
worst area	569.0	880.583128	569.356993	185.200000	515.300000	686.500000	1084.000000	4254.00000
worst smoothness	569.0	0.132369	0.022832	0.071170	0.116600	0.131300	0.146000	0.22260
worst compactness	569.0	0.254265	0.157336	0.027290	0.147200	0.211900	0.339100	1.05800
worst concavity	569.0	0.272188	0.208624	0.000000	0.114500	0.226700	0.382900	1.25200
worst concave points	569.0	0.114606	0.065732	0.000000	0.064930	0.099930	0.161400	0.29100
worst symmetry	569.0	0.290076	0.061867	0.156500	0.250400	0.282200	0.317900	0.66380
worst fractal dimension	569.0	0.083946	0.018061	0.055040	0.071460	0.080040	0.092080	0.20750
target	569.0	0.627417	0.483918	0.000000	0.000000	1.000000	1.000000	1.00000

Checking skewness

In [35]:
numeric_skewness = breast_cancer_df.select_dtypes(include=[np.number]).skew()
print("Skewness:")
print(numeric_skewness)

Skewness:
mean radius 0.942380
mean texture 0.650450
mean perimeter 0.990650
mean area 1.645732
mean smoothness 0.456324
mean compactness 1.190123
mean concavity 1.401180
mean concave points 1.171180
mean symmetry 0.725609
mean fractal dimension 1.304489
radius error 3.088612
texture error 1.646444
perimeter error 3.443615
area error 5.447186
smoothness error 2.314450
compactness error 1.902221
concavity error 5.110463
concave points error 1.444678
symmetry error 2.195133
fractal dimension error 3.923969
worst radius 1.103115
worst texture 0.498321
worst perimeter 1.128164
worst area 1.859373
worst smoothness 0.415426
worst compactness 1.473555
worst concavity 1.150237
worst concave points 0.492616
worst symmetry 1.433928
worst fractal dimension 1.662579
target -0.528461
dtype: float64

Handling the skewness

In [36]:
breast_cancer_new = breast_cancer_df.copy()
skew_threshold = 0.75

skewed_features = numeric_skewness[abs(numeric_skewness) > skew_threshold].index

breast_cancer_new[skewed_features] = np.log1p(breast_cancer_df[skewed_features])
print("Skewness after applying the log1p :\n")
breast_cancer_new.skew()

Skewness after applying the log1p :


```
Out[36]: mean radius      0.348744
mean texture      0.650450
mean perimeter    0.328800
mean area         0.287328
mean smoothness   0.456324
mean compactness  1.043195
mean concavity    1.206982
mean concave points 1.083180
mean symmetry     0.725609
mean fractal dimension 1.272294
radius error      1.713417
texture error     0.629526
perimeter error   1.025375
area error        0.847298
smoothness error  2.283394
compactness error 1.816278
concavity error   4.231738
concave points error 1.399869
symmetry error    2.135874
fractal dimension error 3.874373
worst radius      0.506347
worst texture     0.498321
worst perimeter   0.462758
worst area        0.457614
worst smoothness  0.415426
worst compactness 1.048729
worst concavity   0.717939
worst concave points 0.492616
worst symmetry    1.152984
worst fractal dimension 1.558559
target            -0.528461
dtype: float64
```

Histogram for all columns of Dataframe

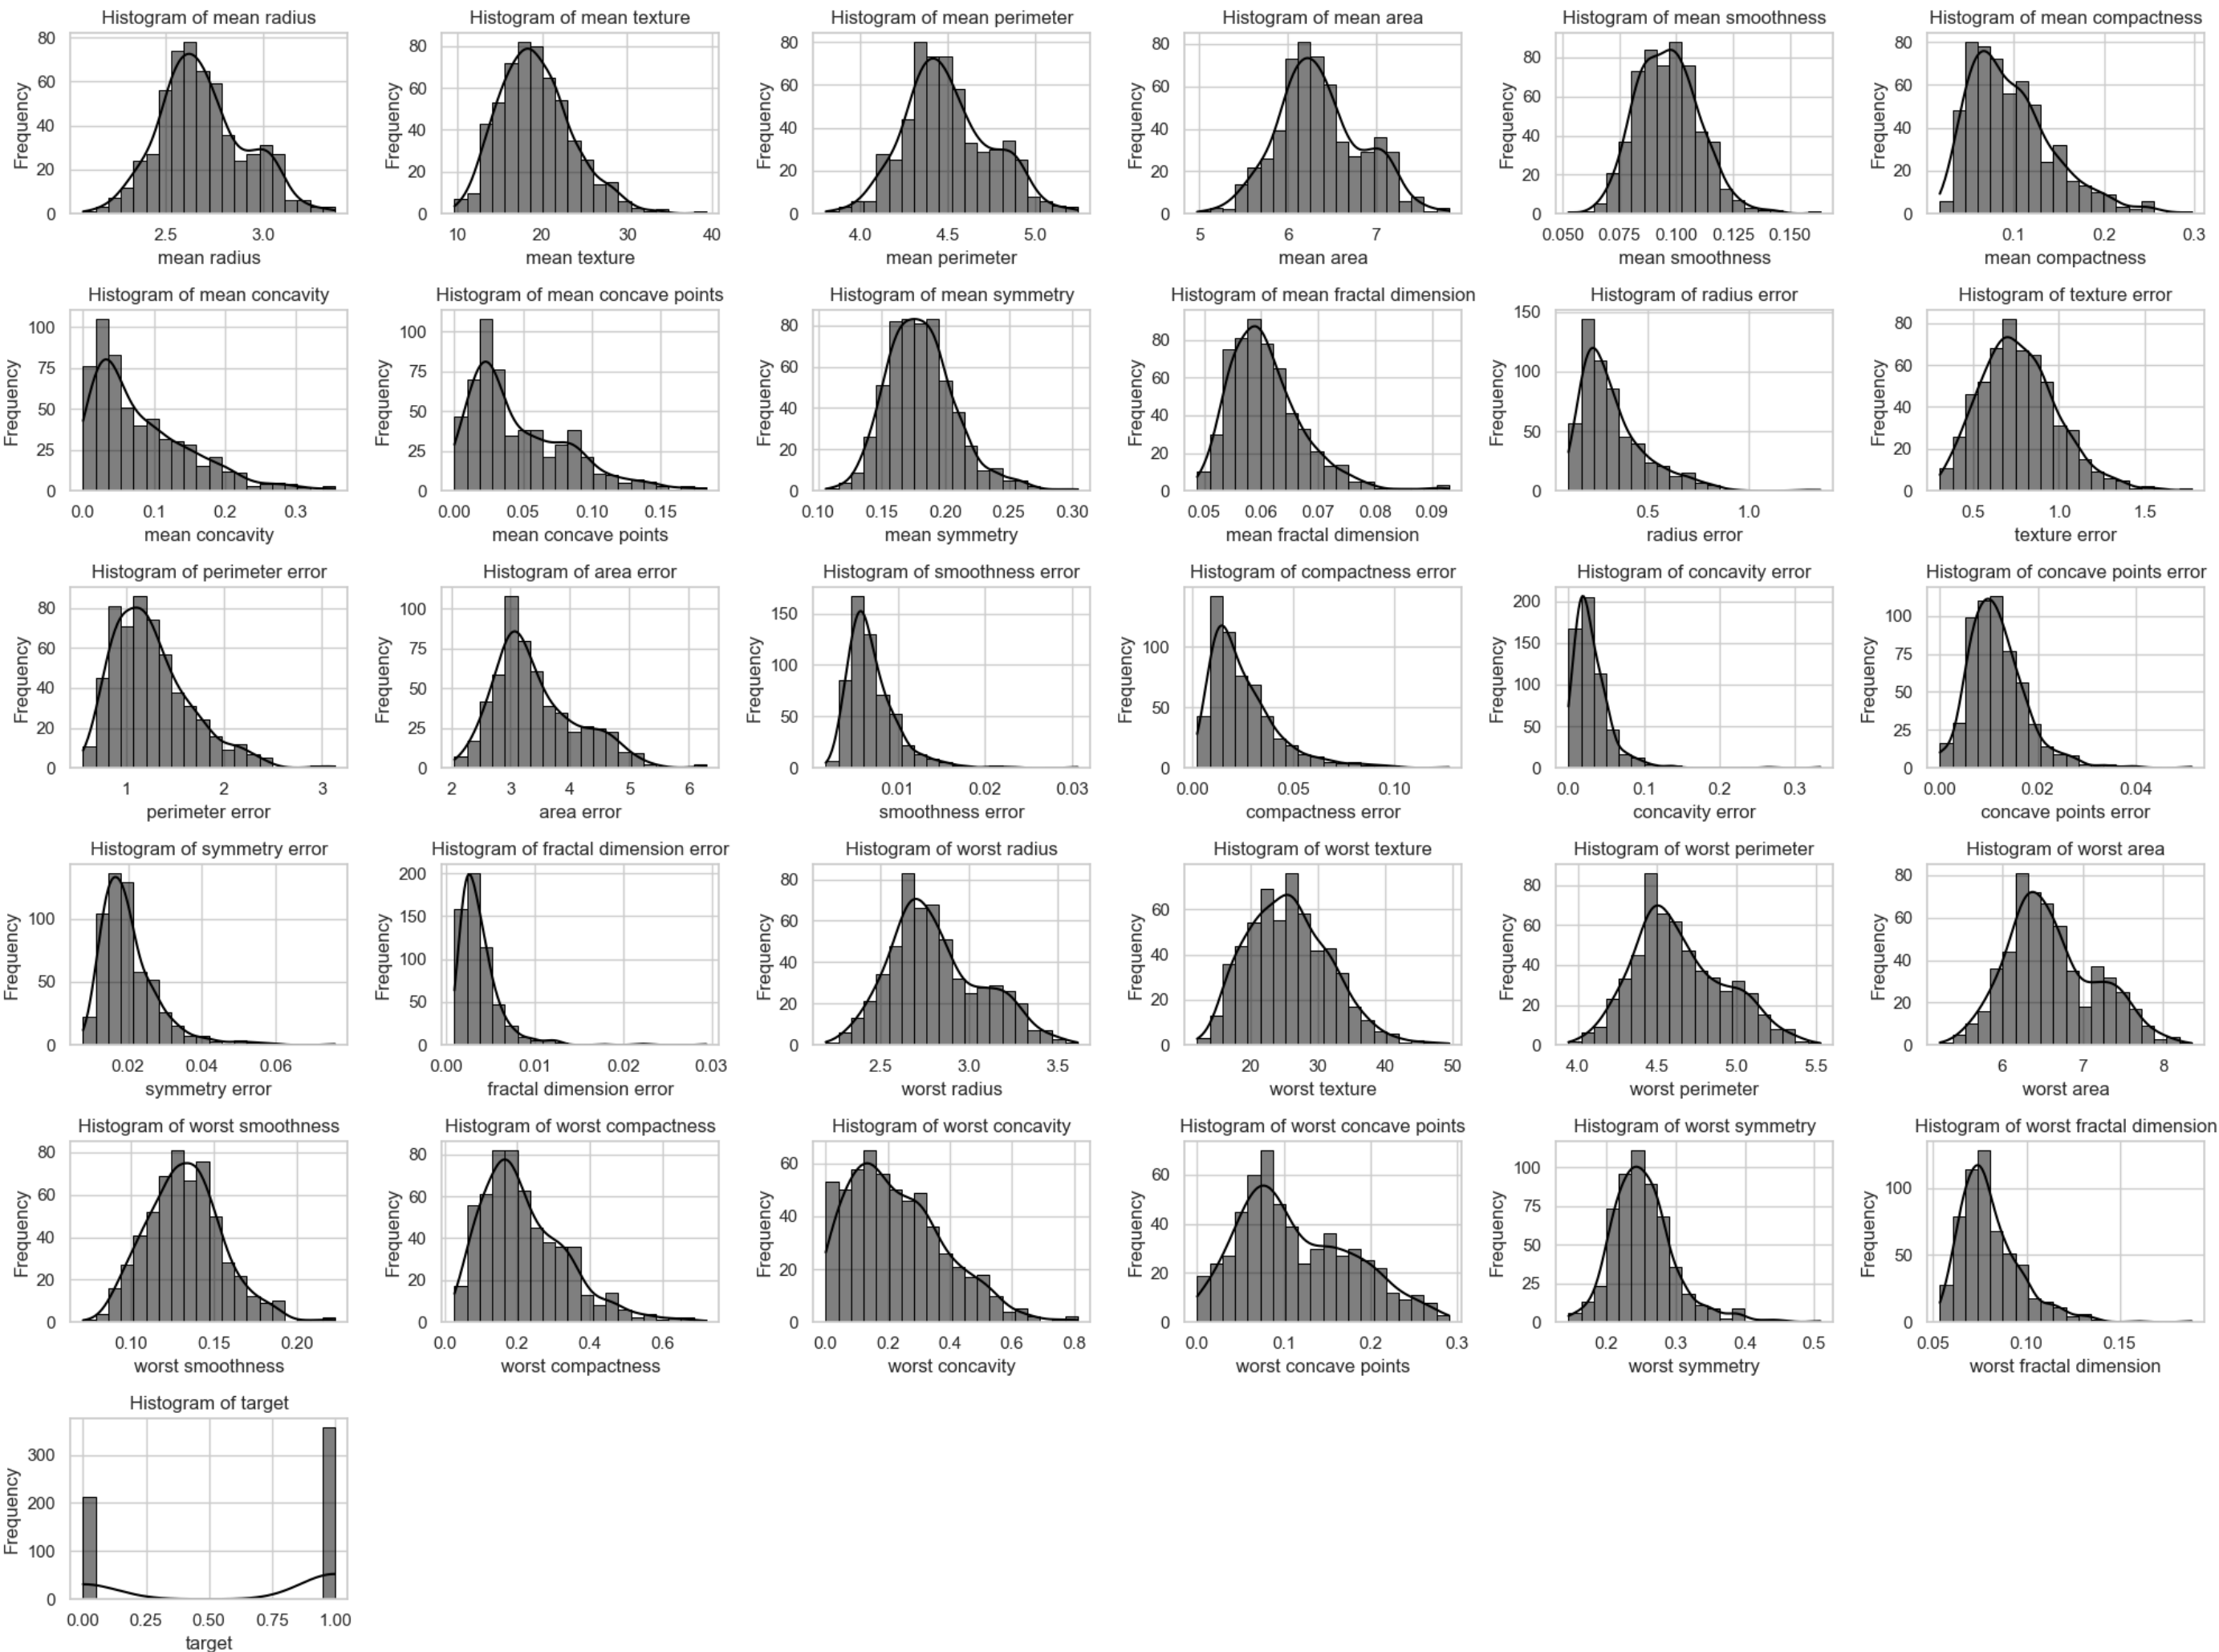
```
In [37]: num_rows = 6
num_cols = 6

fig, axes = plt.subplots(num_rows, num_cols, figsize=(20, 15))

for i, column in enumerate(breast_cancer_new.columns):
    row = i // num_cols
    col = i % num_cols
    ax = axes[row][col]
    sns.histplot(breast_cancer_new[column], kde=True ,bins=20, color='black', edgecolor='black', ax=ax)
    ax.set_xlabel(column)
    ax.set_ylabel('Frequency')
    ax.set_title(f'Histogram of {column}')
    ax.grid(True)

for i in range(len(breast_cancer_new.columns), num_rows * num_cols):
    row = i // num_cols
    col = i % num_cols
    fig.delaxes(axes[row][col])

plt.tight_layout()
plt.show()
```



Box-plot

```
In [86]: num_rows = 6
num_cols = 6


fig, axes = plt.subplots(num_rows, num_cols, figsize=(20, 15))
```



```
for i, column in enumerate(breast_cancer_new.columns):
    row = i // num_cols
    col = i % num_cols
    ax = axes[row][col]
    sns.boxplot(x=breast_cancer_new[column], color='pink', ax=ax)
    ax.set_xlabel(column)
    ax.set_ylabel('Value')
    ax.set_title(f'Boxplot of {column}')
    ax.grid(True)

for i in range(len(breast_cancer_new.columns), num_rows * num_cols):
    row = i // num_cols
    col = i % num_cols
    fig.delaxes(axes[row][col])

plt.tight_layout()
plt.show()
```



Outliers detection using quantile method

```
In [39]: Q1 = breast_cancer_new.quantile(0.25)
Q3 = breast_cancer_new.quantile(0.75)

IQR = Q3 - Q1

threshold = 1.5

outliers = (breast_cancer_new < (Q1 - threshold * IQR)) | (breast_cancer_new > (Q3 + threshold * IQR))

outliers_count = outliers.sum()
total_outliers1 = outliers_count.sum()
print("\nTotal number of outliers:", total_outliers1, '\n')

Total number of outliers: 388
```

Removing outliers using quantile method

```
In [40]: quantile_data = breast_cancer_new[~outliers.any(axis=1)]

print("Shape before removing outliers:", breast_cancer_new.shape)
print("Shape after removing outliers:", quantile_data.shape)

Q2 = quantile_data.quantile(0.25)
Q4 = quantile_data.quantile(0.75)

IQR = Q4 - Q2

threshold = 1.5

outliers2 = (quantile_data < (Q2 - threshold * IQR)) | (quantile_data > (Q4 + threshold * IQR))

print("Outliers:")
outliers_count2 = outliers2.sum()
total_outliers = outliers_count2.sum()
print("\nTotal number of outliers after performing the outlier removal:", total_outliers)

Shape before removing outliers: (569, 31)
Shape after removing outliers: (433, 31)
Outliers:

Total number of outliers after performing the outlier removal: 122

In [41]: print("Date Loss after removing the outliers : ", ((569-433)/569)*100, '%')
```

Date Loss after removing the outliers : 23.901581722319857 %

Box-plot after removing the outliers


```
In [80]: num_rows = 6
num_cols = 6

fig, axes = plt.subplots(num_rows, num_cols, figsize=(20, 15))

for i, column in enumerate(quantile_data.columns):
    row = i // num_cols
    col = i % num_cols
    ax = axes[row][col]
    sns.boxplot(x=quantile_data[column], color='gray', ax=ax)
    ax.set_xlabel(column)
    ax.set_ylabel('Value')
    ax.set_title(f'Boxplot of {column}')
    ax.grid(True)

for i in range(len(quantile_data.columns), num_rows * num_cols):
    row = i // num_cols
    col = i % num_cols
    fig.delaxes(axes[row][col])

plt.tight_layout()
plt.show()
```



Define X and Y

```
In [43]: X = quantile_data.drop(['target'], axis=1)
Y = quantile_data['target']

X = pd.get_dummies(X)
X
```

Out[43]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	radius error	texture error	perimeter error	area error	smoothness error	compactness error	concavity error	concave points error	symm error
1	3.071303	17.77	4.897093	7.190676	0.08474	0.075701	0.083330	0.067818	0.1812	0.055122	0.434053	0.550373	1.481150	4.318554	0.005211	0.012995	0.018429	0.013311	0.018429
2	3.029650	21.25	4.875197	7.093405	0.10960	0.148334	0.180153	0.120357	0.2069	0.058259	0.557098	0.580482	1.720084	4.554193	0.006131	0.039278	0.037604	0.020371	0.020371
4	3.058237	14.34	4.913390	7.168580	0.10030	0.124692	0.180653	0.099212	0.1809	0.057165	0.563722	0.577343	1.862218	4.558498	0.011424	0.024312	0.055321	0.018675	0.018675
6	2.957511	19.98	4.792479	6.947937	0.09463	0.103459	0.106789	0.071390	0.1794	0.055832	0.369285	0.572786	1.430311	4.005695	0.004305	0.013725	0.022290	0.010336	0.010336
7	2.688528	20.83	4.513055	6.361130	0.11890	0.152292	0.089530	0.058127	0.2196	0.071865	0.459638	0.865839	1.580215	3.950474	0.008766	0.029840	0.024576	0.014376	0.014376
...
555	2.423917	27.61	4.199755	5.775793	0.09030	0.073789	0.058259	0.027012	0.1593	0.059466	0.198769	1.175265	0.890768	2.738256	0.011978	0.026992	0.046922	0.017064	0.017064
558	2.746630	22.68	4.578724	6.489357	0.08473	0.124869	0.097943	0.036679	0.1454	0.059655	0.203267	0.745740	1.170623	3.022374	0.004233	0.045346	0.063707	0.015932	0.015932
560	2.711378	27.15	4.525911	6.399260	0.09929	0.106700	0.043653	0.042140	0.1537	0.059881	0.310788	0.913086	1.357895	3.428813	0.007230	0.026428	0.020498	0.016129	0.020498
565	3.050694	28.25	4.884316	7.140453	0.09780	0.098396	0.134531	0.093408	0.1752	0.053854	0.568434	1.242135	1.825033	4.605570	0.005752	0.023941	0.038740	0.016641	0.016641
566	2.867899	28.08	4.694096	6.755885	0.08455	0.097399	0.088478	0.051662	0.1590	0.054943	0.375968	0.729961	1.487270	3.902982	0.005886	0.036631	0.046215	0.015450	0.015450

433 rows × 30 columns

Applying Standard Scaler

```
In [44]: scaler = StandardScaler()
X_scaled = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
print("\nScaled DataFrame using StandardScaler:")
X_scaled
```

Out[44]:

Scaled DataFrame using StandardScaler:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	radius error	texture error	perimeter error	area error	smoothness error	compactness error	concavity error	concave points error
0	2.092937	-0.244628	1.982150	2.089041	-0.754191	-0.289086	0.380612	1.002205	0.294948	-0.891428	1.246656	-0.889997	1.019126	1.736042	-0.589733	-0.669121	-0.423463	0.738427
1	1.870551	0.639468	1.877112	1.849882	1.333734	1.871704	2.315526	2.851326	1.501153	-0.207284	2.277146	-0.739206	1.785724	2.133038	-0.121606	1.816986	0.885998	2.456474
2	2.023178	-1.116021	2.060327	2.034714	0.552652	1.168386	2.325537	2.107108	0.280867	-0.446078	2.332615	-0.754926	2.241746	2.140292	2.572491	0.401331	2.095909	2.043623
3	1.485397	0.316824	1.480303	1.492224	0.076444	0.536692	0.849435	1.127937	0.210466	-0.736689	0.704238	-0.777751	0.856015	1.208947	-1.051196	-0.600054	-0.159821	0.014563
4	0.049287	0.532767	0.139868	0.049454	2.114817	1.989453	0.504518	0.661166	2.097216	2.759851	1.460927	0.689908	1.336967	1.115911	1.219657	0.924245	-0.003723	0.997628
...
428	-1.363473	2.255229	-1.363071	-1.389701	-0.287221	-0.345957	-0.120391	-0.433931	-0.732909	0.055910	-0.723810	2.239561	-0.875058	-0.926398	2.854194	0.654864	1.522310	1.651603
429	0.359496	1.002760	0.454889	0.364723	-0.755031	1.173637	0.672648	-0.093700	-1.385292	0.097006	-0.686136	0.088430	0.022830	-0.447724	-1.087685	2.390929	2.668576	1.376333
430	0.171286	2.138366	0.201538	0.143204	0.467825	0.633108	-0.412284	0.098480	-0.995739	0.146310	0.214334	0.926526	0.623674	0.237031	0.437557	0.601448	-0.282147	1.424229
431	1.982902	2.417821	1.920856	1.965559	0.342684	0.386089	1.403821	1.902864	0.013343	-1.168167	2.372080	2.574460	2.122441	2.219598	-0.314372	0.366244	0.963563	1.548712
432	1.006955	2.374633	1.008348	1.020031	-0.770148	0.356416	0.483494	0.433627	-0.746989	-0.930646	0.760203	0.009408	1.038762	1.035898	-0.246567	1.566554	1.474074	1.258950

433 rows x 30 columns

2.2 Perform PCA (2 components) and LDA (1 components) over the dataset. (20 Points)

Perform PCA

```
In [45]: pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

print("Explained variance ratio (PCA) components:")
print("PCA Component 1:", pca.explained_variance_ratio_[0])
print("PCA Component 2:", pca.explained_variance_ratio_[1])

Explained variance ratio (PCA) components:
PCA Component 1: 0.45393920463471377
PCA Component 2: 0.17305258074751867

In [46]: print(f"Variance explained by the first 2 components: {sum(pca.explained_variance_ratio_[:2])}")

Variance explained by the first 2 components: 0.6269917853822324
```

Perform LDA

```
In [47]: lda = LDA(n_components=1)
X_lda = lda.fit_transform(X_scaled, Y)

print("Explained variance ratio (LDA):", lda.explained_variance_ratio_[0])


Explained variance ratio (LDA): 1.0
```

2.3 Visualise the components and see if its able to segregate the class label in breast_cancer dataset. (10 Points)

Scatter Plot for the PCA (2 components)

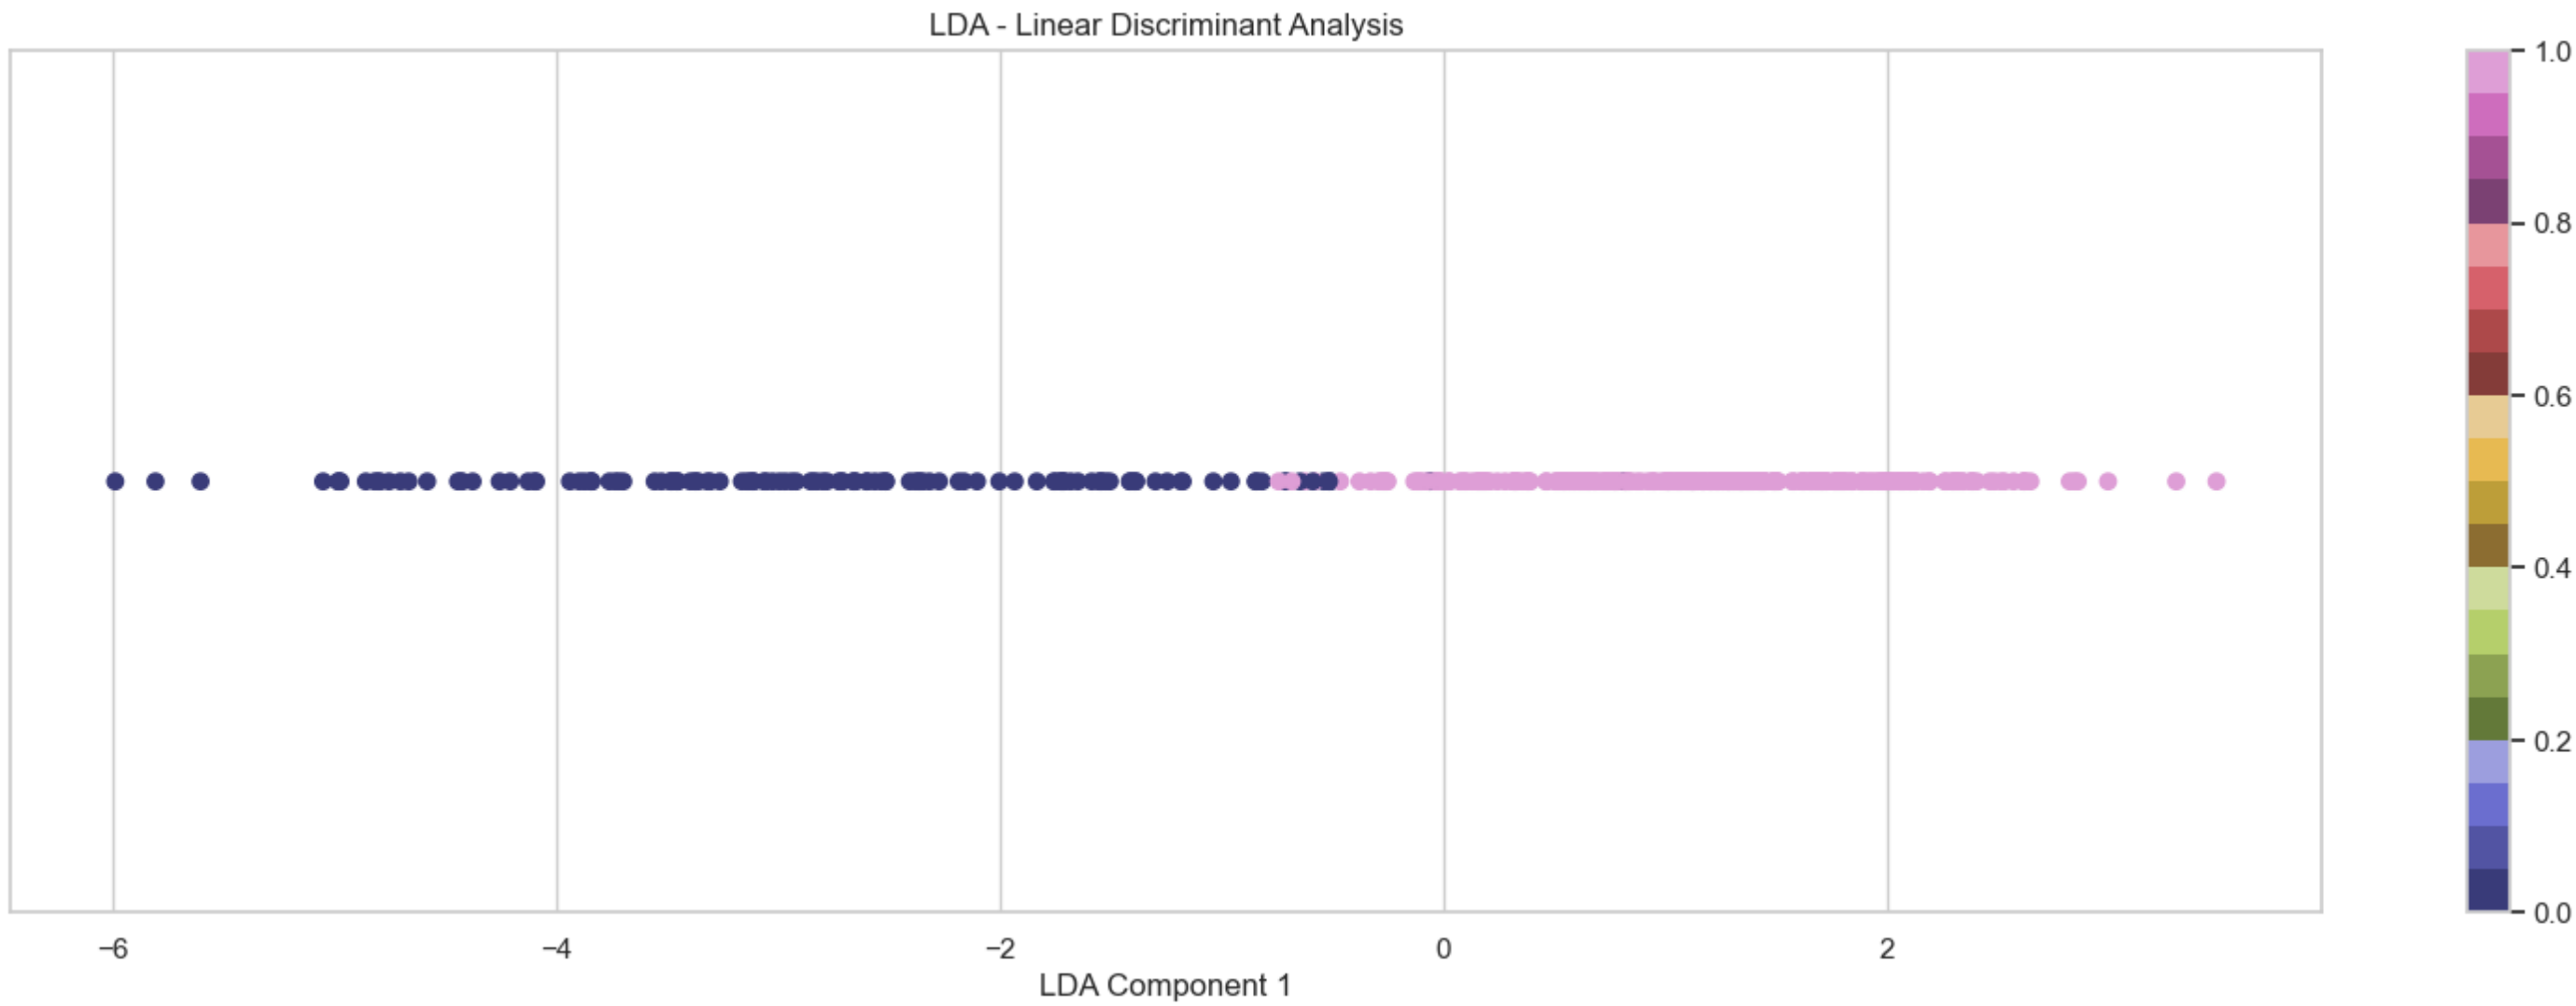
```
In [70]: plt.figure(figsize=(20, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=Y, cmap='Dark2')
plt.xlabel('First principal component')
plt.ylabel('Second principal component')
plt.title('PCA - First two components')
plt.colorbar()

Out[70]: <matplotlib.colorbar.Colorbar at 0x281f009d0>
```



Scatter Plot for LDA (1 components)

```
In [79]: plt.figure(figsize=(20, 6))
plt.scatter(X_lda, [0] * len(X_lda), c=Y, cmap='tab20b')
plt.yticks([])
plt.xlabel('LDA Component 1')
plt.title('LDA - Linear Discriminant Analysis')
plt.colorbar()
plt.show()
```

2.4 What is the maximum variance explained by both the components in PCA and LDA. (10 Points)

Maximum variance explained by components in PCA

```
In [50]: max_variance_pca = np.max(pca.explained_variance_ratio_)
print("Maximum variance explained by PCA components:", max_variance_pca)
```

Maximum variance explained by PCA components: 0.45393920463471377

Maximum variance explained by component in LDA

```
In [51]: variance_lda = lda.explained_variance_ratio_[0]
print("Variance explained by the LDA component:", variance_lda)
```

Variance explained by the LDA component: 1.0

2.5 Comment on the working of PCA and LDA and which one is better for breast_cancer dataset. (5 Points)

Principal Component Analysis (PCA) and Linear Discriminant Analysis (LDA) are two common dimensionality reduction approaches in machine learning and statistics, but they provide different functions and operate in different ways.

PCA (Principal Component Analysis):

1. Purpose:

- PCA is mostly used for unsupervised dimension reduction and feature extraction.
- PCA seeks to identify the directions (principal components) that maximize variance in the data. It does not take into account class labels.
- PCA is an unsupervised technique, which means it does not employ class information during the computation.
- It is useful for lowering the dimensionality of the feature space while preserving as much variance as possible. This makes it appropriate for tasks where maintaining the overall structure of the data is critical.

2. Working:

- PCA identifies orthogonal axes (principal components) in the feature space, with the first principal component capturing the most variation, the second principal component capturing the second most variance, and so on. Each primary component is a linear combination of the original attributes.

3. Objective:

- The goal of PCA is to project data onto a lower-dimensional subspace while retaining as much original variation as possible.

LDA (Linear Discriminant Analysis):

1. Purpose:

- LDA aims to find the directions that maximize the separation between multiple classes in the data.
- LDA is a supervised technique as it takes class labels into account during computation.
- It's particularly useful for classification tasks, as it tries to find the feature subspace that optimally separates different classes.

2. Working:

- LDA aims to maximise the ratio of between-class to within-class variation. It projects the data onto a lower-dimensional space in such a way that the distances between the means of various classes are increased while the variations within each class are reduced.

3. Objective:

- The goal of LDA is to identify the linear combinations of features that best distinguish the classes in the dataset.

Which one is better for the breast_cancer dataset:

- For the breast_cancer dataset, which is a classification problem, LDA may be more appropriate. This is because LDA directly models class differences and seeks to identify the feature subspace that maximizes class separation. Because the dataset includes class labels, LDA can use this information to identify the most discriminative features.
- If computing economy is an issue, PCA can also be used for visualization or to reduce the dimensionality of a dataset.
- Finally, the decision between PCA and LDA is determined by the precise analytic aims and the type of the dataset.
 1. If classification is the goal and maximum class separation is important, LDA may be preferable.
 2. If the goal is dimensionality reduction or visualization, PCA.